

# Structuring a DS Project

Lecture 3, DSC 180A

# Announcements


Assignment 1 Comments: Comment on Grading; Code Review Friday.

Working on projects:

- Work a little every day (things always go wrong!)
- Work in parallel on multiple pieces
- Don't let a single step be a 'blocker'.

# Data Ingestion


- Task: get data from internet to computer.
- Make it easy to (incrementally) change data ingested.
- Rerun to 'refresh' with new data.
- Run data pull on different servers to reproduce results



How to organize code to achieve this?

# Data Ingestion: `etl.py`

- Library code: contains functions for importing by other processes.
- Good for interactive use; reusable.
- Written as generically as practicable.
- Contains logic not necessary for a *consumer* of the project to know.
- Contains data collection logic other *developers* might want to expand on, when forking a project.
- Library code know nothing of what calls it!



Project

- README.md
- data-params.json
- etl.py
- run.py

# Data Ingestion: `etl.py`

- Library code: contains functions for importing by other processes.
- Written as generically as practicable.
- Good for interactive use; reusable.
- Contains logic not necessary for a *consumer* of the project to know.
- Contains data collection logic other *developers* might want to expand on, when forking a project.
- Library code know nothing of what calls it!

```
'''
etl.py contains functions used to download tables for different
teams and years.
'''


def get_season(team, year):
    '''
    return a table of season statistics for a
    given team and year.
    '''
    ...
    return ...

def get_data(years, teams, outdir):
    '''
    downloads and saves tables at the specified output
    directory
    for the given years and teams.

    :param: years: a list of seasons to collect
    :param: teams: a list of teams to collect
    :param: outpath: the directory to which to save the data.
    '''
    ...
    return
```

# Data Ingestion: `data-params.json`

- Configuration: parameters for different investigations and experiments.
- E.g. Parameterize across time/space.
- Used by the *consumer* of the project. Shouldn't require a knowledge of the source code!
- Helps log the results of different experiments.



Project

- README.md
- data-params.json
- etl.py
- run.py


# Data Ingestion: data-params.json

- Configuration: parameters for different investigations and experiments.
- E.g. Parameterize across time/space.
- Used by the *consumer* of the project. Shouldn't require a knowledge of the source code!
- Helps log the results of different experiments.

```
{  
  "years": [2015, 2016, 2017, 2018, 2019],  
  "teams": ["sfo", "gnb"],  
  "outpath": "data/raw"  
}
```

# Data Ingestion: `run.py`

- Script: Builds (common portions of) the project.
- Imports and runs library code: gives examples of *code usage*.
- Current: hand-made `run.py`
- Also possible to use specialized tools: python CLI (e.g. argparse), Makefiles, Maven, etc...



Project

- ├── README.md
- ├── data-params.json
- ├── etl.py
- └── run.py



# Data Ingestion: `run.py`

- Imports library code (`get_data`)
- Run as a script:
  - `python run.py data`
- Shebang: `#!/usr/bin/env python`
  - Specifies which python interpreter to use.
- `main` function strings together library functions with parameters in config.
- `__name__ == '__main__'...` returns true only when file is run as a script. Should only have minimal code inside.

```
#!/usr/bin/env python

import sys
import json

from etl import get_data

def main(targets):

    if 'data' in targets:
        with open('data-params.json') as fh:
            data_cfg = json.load(fh)

            # make the data target
            get_data(**data_cfg)

    return

if __name__ == '__main__':
    targets = sys.argv[1:]
    main(targets)
```

# Project Structure:

- As project grows, so does code complexity!
- As a project grows, it becomes unclear:
  - how code should be run...
  - what the code *does*...
  - if the code is correct...

```
In [10]: import numpy as np
import pandas as pd

In [11]: import seaborn as sns

In [12]: flights = sns.load_dataset('flights')
flights = flights.reset_index()
flights['log(passengers)'] = np.log(flights.passengers)
flights['sqrt(passengers)'] = np.sqrt(flights.passengers)

In [13]: print(linregress(flights['index'], flights['passengers']).rvalue ** 2)
# sns.regplot(data=flights, x='index', y='passengers')

NameError                                Traceback (most recent call last)
<ipython-input-13-082f0304aa1a> in <module>()
----> 1 print(linregress(flights['index'], flights['passengers']).rvalue ** 2)
      2
      3 # sns.regplot(data=flights, x='index', y='passengers')

NameError: name 'linregress' is not defined

In [ ]: print(linregress(flights['index'], flights['log(passengers)']).rvalue ** 2)
# sns.regplot(data=flights.reset_index(), x='index', y='log(passengers)')

In [ ]: flights[flights.dtypes != np.number].index

In [ ]: flights.dtypes

In [ ]:

In [ ]: data = (
pd.concat([diamonds.copy(), create_features(diamonds)], axis=1)
.drop('price', axis=1)
# .drop('carat', axis=1)
.select dtypes={'number':})

In [ ]: from sklearn.linear_model import LinearRegression
from scipy.stats import linregress

In [ ]: y = diamonds.price
res = {}
for x in data.columns:
    X = data[[x]]

    lr = LinearRegression()
    lr.fit(X, y)
    res[x] = lr.score(X, y)
```

# Why Care About Project Structure?

*We're not talking about bikeshedding the indentation aesthetics or pedantic formatting standards — ultimately, data science code quality is about correctness and reproducibility.*

## Cookie Cutter Data Science

- Clear and consistent project organization encourages software development best-practices and readable code.
- Such habits yield more consistently correct code that's more easily fixed and adapted to other tasks.

# Why Care About Project Structure?

*We're not talking about bikeshedding the indentation aesthetics or pedantic formatting standards — ultimately, data science code quality is about correctness and reproducibility.*

## Cookie Cutter Data Science

- Clear and consistent project organization encourages software development best-practices and readable code.
- Such habits yield more consistently correct code that's more easily fixed and adapted to other tasks.
- We will follow the opinions of [Cookie Cutter Data Science](#)

# An *Example* Project Template

```
|— .gitignore      <- files to keep out of version control (e.g. data/binaries)
|— run.py         <- run.py with commands like `make data` or `make train`
|— README.md      <- The top-level README for developers using this project.
|— data
|   |— temp       <- Intermediate data that has been transformed.
|   |— out        <- The final, canonical data sets for modeling.
|   |— raw        <- The original, immutable data dump.
|— notebooks      <- Jupyter notebooks (presentation only).
|— references      <- Data dictionaries, explanatory materials.
|— requirements.txt <- For reproducing the analysis environment, e.g.
|                   generated with `pip freeze > requirements.txt`
|— src            <- Source code for use in this project.
|   |— data       <- Scripts to download or generate data
|   |   |— make dataset.py
|   |— features   <- Scripts to turn raw data into features for modeling
|   |   |— build features.py
|   |— models     <- Scripts to train models and make predictions
|   |   |— predict model.py
|   |   |— train model.py
|   |— visualization <- Scripts to create exploratory and results oriented viz
|       |— visualize.py
```

# Results are Derived from Immutable Raw Data

- Data is immutable: *never* edit raw data
  - Raw data is always (re)ingested from elsewhere.
  - File-path may be a *symbolic link* (shortcut), if stored locally.
  - Raw data never changes => doesn't need version control! (.gitignore)
- Final data is always reproducible from raw data (with run.py)
- Temp data holds data 'useful to keep around' for development, analysis, debugging, etc...

```
...
├── data
│   ├── temp      <- Intermediate data that has been transformed.
│   ├── out       <- The final, canonical data sets for modeling.
│   └── raw       <- The original, immutable data dump.
...
```

# Notebooks are for Analysis and Communication

- Notebooks are great for communication, analysis, and initial development.
  - Use to create up-to-date, reproducible, static HTML reports.
- Complicated code in notebooks are hard to understand and don't work well with version control and collaboration.
- Notebooks should:
  - Mostly call library functions in src, with very simple code logic.
  - Never “copy-paste” code between notebooks -- if it's reusable, put it in a library function.
  - Name it something descriptive: `03-fraenkel-prelim-EDA.ipynb`

```
...  
└─ notebooks          <- Jupyter notebooks (presentation only).  
...
```

# Build from the Environment Up

- To reproduce a project from scratch, must also reproduce the computational environment on which it was run.
- requirements.txt contains all python libraries needed for running the project.
- `git clone project => mk virtualenv => pip install requirements.txt`
- When a project has more complicated requirements, may need to use container approach (e.g. Docker or Vagrant).

```
...
|— requirements.txt  <- For reproducing the analysis environment, e.g.
|                   generated with `pip freeze > requirements.txt`
...
```