

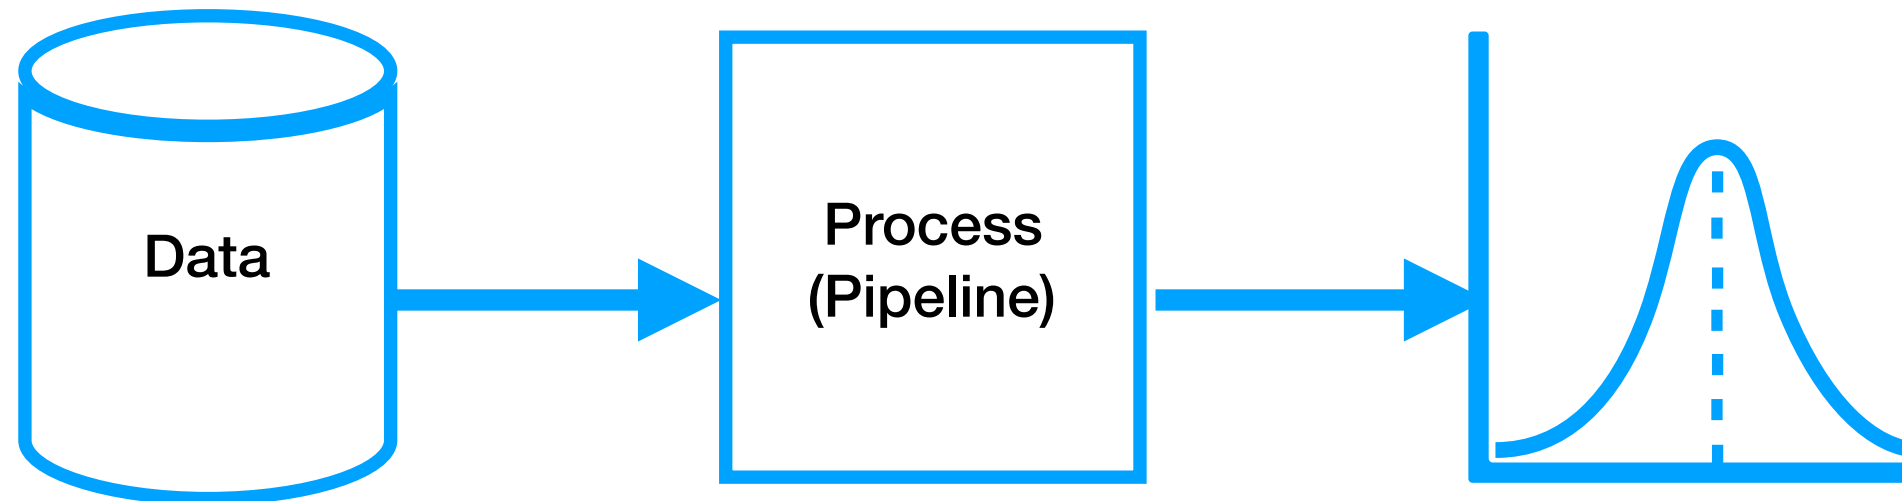
# Workflow / Patterns

Lecture 05

# Outline

- Writing for pipeline evolution
- Practical patterns for big data
- Interacting with the file-system

# The End Result




- The clean abstraction of a finished pipeline.
- Reality is an evolving process:
  - Investigation: simple => grown in scope and sophistication.
  - Data becomes multi-faceted (sources; time; space)
  - Analyses become refined

# Example: Football


- Topic: Football players and injuries
- Data: Injured Reserve (IR) list & Football Statistics
- Narrow Question: Do injuries affect team success?
- Not to be forgotten: Broader social context!

# Do Injuries Affect Teams?



- Player level model? Team Level?
- Even answering basic questions requires iteration:
  - Requires breaking up the pipeline; collecting data on disk

# Pipeline, Refined.




- Persist data locally to save processing and cost, ease debugging.
- Configured to *also* run w/o persistence (e.g post-exploration).

# Considerations when Designing a Pipeline

- Further questions: What causes player injury? What rule changes would have largest affect? New technologies?
- Possible new data sources: player-level data; play-by-play data; news stories; game video.
- For each of these datasets:
  - Small data? (Yes? Don't bother optimizing!)
  - Heterogeneous or large? (Yes? Don't join prematurely!)

# (Future) Pipeline



- Each step is an abstraction hiding complicated processing



# Possible Pattern

```
# Data Ingestion Funcs: may write to disk or return in Mem
data_ing = {'IR': data_ing_IR, 'Team': data_ing_Team, ...}

# Data Processing Funcs: may read from disk, or in mem objs
procs = {'IR': proc_IR, 'Team': proc_Team, ...}

joined = ...

for d in datasets:
    ing_func, proc_func = data_ing['d'], data_proc['d']
    dataset = proc_func(ing_func(...))
    joined = joined.merge(dataset)

out = mdl(joined)
```

# Patterns: what and where

- Data ingestion code can *either* write to disk or output DF
- Processing code can *either* read from disk, or input DF
- 'Driver functions' build the pipeline from small functions:
  - Implement logic for persisting/reading to/from disk.
  - Creates a *directed acyclic processing graph (DAG)*:
    - *Nodes are (data) sources and targets*
    - *between edges that process the data.*

# Patterns: How

- Processing DAGs help organize and structure the code.
- The processing pipeline hides the complexity of code *up close* and manages it as the code grows.
- As code and data grow, need patterns to manage efficiency.

# Memory and Speed

- Need ways to manage memory and speed constraints, without too much compromise on code generality.
- The *streaming vs batch* processing paradigm:
  - Streaming: processing data as it comes requires a small memory footprint.
  - Batch: processing data in larger chunks allows for efficient memory allocation (fast).
- Reality: a spectrum (from batch-size 1 to  $\infty$ ).

# Streaming in Python

Iterators *stream* through elements of an iterable; only the current element is stored in memory.

```
def natural_numbers():
    k = 0
    while True:
        yield k
        k += 1

N = natural_numbers()

for i in N:
    if (i**13 + 122) % 57 == 0:
        break # if no break, wouldn't finish

print(i)
```

49

Streaming through the natural numbers

# Examples: Streaming

- Iterating through bytes, integers, lines...
- wget: download *buffered byte streams* (e.g. 256K buffers)
- Iterating through larger 'chunks' or 'batches':
  - Choose a csv chunksize? Choose a convenient 'unit' of processing? (year, page, person, observation-group)
- When data comfortably fits in memory, a batch can contain all the data.

# Processing & Streaming

- Write library functions that:
  - inputs a chunk & outputs a processed chunk.
- Create a 'driver function' that:
  - regulates (the size of) data intake, and
  - compose library functions, feeding one chunk at a time.
- May involve:
  - Writing many files to disk,
  - Collecting processed chunks into appropriate 'units'

# Streaming huge files

- Question: what should a chunk be for course domains? (i.e. what is the unit to iterate through?)
- If non-tabular organize many files on disk (~partitions)
  - Store in many files for possible parallelism advantages
- If tabular, make a relational database (e.g. local sqlite)